JYOTHIR RAGHAVALU BHOGI

Email: bhogiyothir.dev@gmail.com
Mobile: +91 6281073800

Jyothir-dev | LinkedIn
Jyothir-dev (Jyothir) (github.com)

EDUCATION

Lovely Professional University — Jalandhar, India
Bachelor of Technology (Computer Science & Engineering) — **July 2022 - May 2026**

SKILLS SUMMARY

- **Programming**: Python, SQL
- **Data Infrastructure**: Scalable Data Architecture, Data Lakes, Data Pipelines, Distributed Systems, Query Engines, Large-Scale Infrastructure
- **Data Processing**: Apache Kafka, Apache Spark, Apache Airflow, Trino, Apache Iceberg
- **Databases**: MySQL, Redis, Query Optimization, Indexing
- **Experimentation**: A/B Testing Infrastructure, Experiment Design
- **Performance**: Cost Optimization, Performance Tuning, Reliability Engineering, Observability
- **Backend**: REST APIs, System Design
- **Collaboration**: Cross-Functional Collaboration, Problem Solving, Ownership

WORK EXPERIENCE

DATA ENGINEER INTERN | flickDONE | LINK — **June 25- August 25**

- Optimized Python-based data pipelines processing 50,000+ daily records, improving reliability and throughput by 30%.
- Engineered scalable SQL ingestion workflows with indexing, reducing query latency and improving performance efficiency.
- Developed Kafka-integrated real-time data pipelines supporting distributed systems and scalable query engines.
- Collaborated cross-functionally deploying production-grade REST APIs aligned with SDLC reliability standards.
- Implemented structured logging and monitoring dashboards, resolving incidents and reducing processing overhead.

PROJECTS

Generative AI Product Recommender System | LINK — **June 25- August 2025**

- Built scalable Python-based recommendation engine using SQL, distributed systems, and vector retrieval pipelines.
- Implemented A/B testing infrastructure evaluating ranking models, improving recommendation relevance by 22%.
- Designed low-latency data pipelines integrating embeddings and query engine optimizations for real-time retrieval.
- Optimized data storage workflows reducing infrastructure costs by 18% through efficient indexing strategies.

Rails TODO App | LINK — **September 24- November 2024**

- Developed backend services using MVC architecture with SQL database integration and modular REST APIs.
- Optimized MySQL queries and Redis caching mechanisms reducing response latency and improving scalability.
- Implemented structured testing, performance monitoring, and deployment workflows ensuring reliable production releases.
- Ensured clean system design principles aligned with scalable infrastructure and maintainable code standards.

Morning Market Brief | LINK — **June 24- August 2024**

- Developed Python-based ETL pipelines processing financial datasets using SQL-driven data transformations.
- Built scalable data lake architecture enabling structured ingestion, validation, and integration with distributed query engines.
- Implemented REST APIs and automation workflows improving reporting accuracy and system reliability by 35%.
- Optimized storage and processing costs through indexing, performance tuning, and infrastructure monitoring practices.

RESEARCH AND PUBLICATIONS

Multi-Stage Ensemble for Driver Fatigue Recognition | IEEE, 2025

- Developed PyTorch-based ensemble pipeline improving classification accuracy and reducing false positives by 20%.
- Engineered feature extraction workflows leveraging embeddings and scalable preprocessing for distributed training.
- Applied structured experimentation, hyperparameter tuning, and validation pipelines ensuring reproducible performance.
- Optimized computational performance reducing processing time and improving system scalability for large datasets.

CERTIFICATES

Python, Data Science & ML (Cipherschools) | CERTIFICATE — **August 2024**

- Mastered Python programming, SQL fundamentals, data structures, and scalable backend system development principles.
- Applied machine learning pipelines, feature engineering, and experimentation workflows on structured datasets.

Generative AI with LLMs (Coursera) | CERTIFICATE — **March 2024**

- Learned transformer architectures, embeddings, distributed inference concepts, and scalable AI infrastructure fundamentals.
- Built LLM-integrated backend systems leveraging REST APIs, experimentation frameworks, and performance optimization.